

# Aloha from Lua

## Table of Contents

<a href="#">Introduction</a> .....	3
<a href="#">eLua</a> .....	3
<a href="#">Building eLua</a> .....	4
<a href="#">Demonstrating eLua</a> .....	7
<a href="#">Multi-Processing with eLua</a> .....	8
<a href="#">Serial ALOHA</a> .....	10
<a href="#">Preparing multiple Propellers</a> .....	10
<a href="#">Building eLua with Serial ALOHA support</a> .....	12
<a href="#">A simple demo</a> .....	15
<a href="#">Loading the simple demo</a> .....	15
<a href="#">Executing the simple demo</a> .....	16
<a href="#">The simple demo files</a> .....	16
<a href="#">client.lua</a> .....	16
<a href="#">server.lua</a> .....	17
<a href="#">remote.lua</a> .....	18
<a href="#">common.lua</a> .....	19
<a href="#">null.lua</a> .....	20
<a href="#">Putting it all together</a> .....	21
<a href="#">WiFi ALOHA</a> .....	22
<a href="#">Preparing multiple Propellers</a> .....	22
<a href="#">Building eLua with WiFi ALOHA support</a> .....	23
<a href="#">A hybrid demo</a> .....	25
<a href="#">Loading the hybrid demo</a> .....	25
<a href="#">Executing the hybrid demo</a> .....	26
<a href="#">A HTTP demo</a> .....	27
<a href="#">Loading the HTTP demo</a> .....	27
<a href="#">Executing the HTTP demo</a> .....	28
<a href="#">Technical Notes</a> .....	30
<a href="#">File name conventions</a> .....	30
<a href="#">Lua initialization module</a> .....	30
<a href="#">Hub RAM Allocation</a> .....	31
<a href="#">Cog allocation</a> .....	32
<a href="#">The binser module</a> .....	33
<a href="#">The Serial ALOHA protocol</a> .....	34
<a href="#">aloha_tx</a> .....	35
<a href="#">aloha_rx</a> .....	36
<a href="#">WiFi/RPC network definition</a> .....	36
<a href="#">The WiFi ALOHA protocol</a> .....	38
<a href="#">The ALOHA dispatcher</a> .....	38

<a href="#"><u>Background tasks</u></a> .....	40
<a href="#"><u>Appendix A - A more complete demo</u></a> .....	42
<a href="#"><u>client.lua</u></a> .....	42
<a href="#"><u>server.lua</u></a> .....	43
<a href="#"><u>common.lua</u></a> .....	44
<a href="#"><u>remote.lua</u></a> .....	45
<a href="#"><u>Appendix B - Life</u></a> .....	46

## Introduction

**Lua** is a C-based embedded<sup>1</sup> scripting language that has found popularity in games and other applications where providing the ability for end-users to customize applications easily and rapidly is highly desirable. Lua is not only written in C, it is intended to be highly *inter-operable* with C – you can call Lua programs from C, and vice-versa. But Lua can also be used stand-alone.

An introduction to Lua on the Propeller 2 is given in the document **Lua on the Propeller 2 with Catalina**. That document also contains a description of Catalina's multi-processing extensions to Lua, which enable Lua to take advantage of multiple cogs on multiple threads.

This document is about a *different* extension to Lua that Catalina offers - **eLua** offers a simple way to execute two instances of Lua on the same Propeller using a simple client/server architecture, which provides another way for Lua to take advantage of the Propeller's multi-processing capabilities.

eLua allows Lua to get around both the speed and space limitations that could otherwise make Lua impractical on a microcontroller. Lua programs can potentially have both *speed* and *space*, as well as the *multi-processor* capabilities offered by the Propeller.

eLua programs can also be easily extended to multiple propellers, using either a very simple serial protocol called **ALOHA**, or by using a WiFi based Remote Procedure Call (RPC) protocol. First, this document will describe eLua by itself, then eLua with ALOHA, then eLua with WiFi.

## eLua

The main purpose of eLua is to enable two Lua instances to execute in parallel. Typically, one Lua instance will execute from Hub RAM at full propeller speed but with limited program and data space, while the other Lua instance will execute from XMM RAM, at slower speeds but with potentially megabytes of code and data space.

The interaction between the two Lua instances is via a very simple client/server architecture, where the server provides one or more Lua functions that the client can call whenever it needs something done that would require too much code space or data space for it to do itself. The server can also run a “background” task when it is not servicing calls from the client, allowing for an easy to understand and easy to program way to build a multi-processing application.

---

<sup>1</sup> The term "embedded" here means embedded within other programs, not embedded in hardware.

## Building eLua

There are several variants of essentially the same eLua client/server program in the folder `demos/elua` - two main ones, plus a few others that may be useful in particular circumstances.

Here is a brief description of the two main **eLua** variants:

- elua** a CMM<sup>2</sup> Lua client with Lua multi-processing support enabled, and an XMM LARGE Lua server with an 8K cache, both *with* a Lua parser.
- eluax** a CMM Lua client with Lua multi-processing support enabled, and an XMM LARGE Lua server with a 64K cache, both *without* a Lua parser.

The other **eLua** variants omit the multi-processor support, which frees up more Hub RAM when these capabilities are not required:

- eluafx** an NMM Lua client and an XMM LARGE Lua server with an 8K cache, both without a Lua parser or multi-processing support. Because the client is a NATIVE program, it will execute faster than the other eLua variants - but there is less Hub RAM available, so only smaller Lua clients (or servers, depending on how the Hub RAM is allocated) can be executed.
- elus** a CMM Lua client and an XMM LARGE Lua server with an 8K cache, both with a Lua parser, but without multi-processing support.
- eluassx** a CMM Lua client and an XMM LARGE Lua server with a 64K cache, both without a Lua parser or multi-processing support.

Compile them all using the **build\_all** script. For example:

**build\_all P2\_EDGE**

or compile them individually by using the following Catapult commands (this example is for the P2\_EDGE<sup>3</sup>):

```
set CATALINA_DEFINE=P2_EDGE
catapult elua.c
catapult eluax.c
catapult elus.c
catapult eluafx.c
catapult eluassx.c
```

The variants differ only in the attributes of the various Catapult pragmas, which tell Catalina what memory model to use for the primary (server) and secondary (client) programs, their stack sizes, address, cache size, etc.

<sup>2</sup> CMM, LMM, NMM and XMM are different memory/execution models for code on the Propeller 2. See the [Catalina Reference Manual \(Propeller 2\)](#) for details.

<sup>3</sup> All the examples in this document assume Windows. On Linux instead of **set** you would say:  
**export CATALINA\_DEFINE=P2\_EDGE**

All the eLua variants should build "as is" when compiled under Windows<sup>4</sup> for a P2 Edge with on-board PSRAM (i.e. the P2-EC32MB), or a P2 Evaluation board with the HyperRAM add-on board. With minor modification - typically, to the address to use for the client program - they can also be built for other Propeller 2 platforms with supported PSRAM. As usual with Catapult programs, the programs will tell you the address to use when they are compiled and executed if this needs to be changed.

All the eLua variants load the client and server Lua files specified on the command line - the first argument specifies the client file, and the second specifies the server file. The variants that include a Lua parser (such as **elua**) can accept either text or compiled binary Lua files (e.g. *client.lua* or *client.lux*). The variants that do *not* include a Lua parser (such as **eluax**) can accept only compiled binary Lua files (e.g. *client.lux*). If no files are specified, the defaults are *client.lua* and *server.lua* for those variants that include a Lua parser, and *client.lux* and *server.lux* for those that do not. All the eLua programs will add the appropriate extension if it is not specified, so for example

**elua client server**

is the same as

**elua client.lua server.lua**

and

**eluax client server**

is the same as

**eluax client.lux server.lux**

All eLua programs use the Catalina Registry for their client/server interaction - this enables the server to offer multiple services, and also adds lock protection so that the services can be used safely in multi-processing applications.

The Catalina registry supports only a limited number of service parameter profiles - basically, those that are required to implement Catalina plugins.

However, the basic "short" service profile can also be used to pass a pointer to the shared data structure as the parameter. This allows for arbitrary data to be exchanged between client and server. In eLua, the shared data is primarily used to store and share the names of the client and server files, and also to synchronize the startup between the client and the server (so that the client does not try to call services provided by the server before the server is ready).

Also, specifically for Lua a "serial" service profile has been added that can be used to accept and return any simple Lua data type. They are passed in a "serialized" format. A binary serialize/deserialize library (called **binser**) is provided for this purpose. Any Lua function can be implemented using the serial service type, because it can accept and return one or more of any of the basic Lua data types - including Lua functions. This

---

<sup>4</sup> All the eLua variants have been configured for compilation on Windows. The address specified in the catapult **secondary** pragma may need to be modified when compiled under Linux. The program will tell you the correct address to use when compiled and executed if it needs to be modified.

means a client can not only exchange arbitrary data with the server - it can also exchange functions with the server for remote execution (the simple demo provided has an example of doing this). The **binser** module is described in more detail in the **Technical Notes** section of this document.

Note that ALOHA protocol currently limits serialized service data to a total size of 2048 bytes. This limit applies to the data after it has been serialized by the **binser** module, so (for example) a single string would have to be less than about 2040 bytes (to allow for the **binser** overhead).

Note that the eLua programs DO NOT generally need to be recompiled just to execute different Lua programs. The main reasons eLua would need to be recompiled would be to adjust the allocation of Hub RAM between the Lua client and the Lua server, to alter plugins used, or to alter the default modules loaded by Lua. See the **Technical Notes** section of this document for more detail on Hub RAM allocation.

## Demonstrating eLua

To demonstrate eLua, after building the binaries copy the following files from *demos/elua* to an SD card containing Catalyst:

<i>elua.bin</i>	the eLua program with a Lua parser
<i>eluax.bin</i>	the eLua program without a Lua parser

Then add the files from the *demos/elua/example* folder:

<i>binser.lua</i>	the <b>binser</b> module (more on this later)
<i>common.lua</i>	common definitions for the clients and servers
<i>client.lua</i>	an example eLua client
<i>server.lua</i>	an example eLua server
<i>remote.lua</i>	an ALOHA proxy client/server (more on this later)
<i>serverbg.lua</i>	an example eLua server with a background task
<i>null.lua</i>	a “null” client or server (does nothing)
<i>rebuild</i>	a Catalyst script to re-compile all the Lua files

There may also be a README.TXT file.

These files implement an example eLua program. The example needs to be compiled before being executed. A catalyst script to recompile all the Lua files is provided called **rebuild**. To execute it use the following Catalyst command:

**exec rebuild**

Since the files have the names that will be loaded by default, to execute the example simply execute any of the eLua variants without specifying any parameters. For example:

**elua**

or

**eluax**

This example implements five different services, intended to illustrate various eLua features:

<b>add</b>	accepts two numbers and returns their sum. Demonstrates that a service can accept multiple arguments.
<b>modrem</b>	accepts two numbers and returns their modulus and remainder. Demonstrate that a service can return multiple results.
<b>invert</b>	accepts a Lua table and returns a Lua table that “inverts” the value of each table element. Demonstrates that a service can accept and return a Lua table.
<b>invoke</b>	accepts a Lua table that contains a number and a Lua function, invokes the function on the number, and returns the result. Demonstrates that a service can accept any simple Lua data type, including Lua functions.

**quit** causes the server to terminate. Provides a mechanism for orderly program termination.

The *common.lua* file contains the definition of all the services required by the *client.lua* program, and *proxy* functions that can be used to call them. The *server.lua* program is slightly more complex, but essentially just adds a *wrapper* function around each service that allows it to be called by the Lua dispatcher. The client doesn't need to know any of these details - it simply calls the proxy functions as if they were local functions and is unaware that they may actually be implemented by the server and not the client.

Note that if you change any of the Lua text files, you should also recompile the corresponding binary versions. For example:

**luac -o client.lux client.lua**

or

**clua client**

Instead of discussing this example in detail, it is better to start with a simpler demo. The simple demo given later in this document functions perfectly well when executed as an eLua example with or without ALOHA. So it is possible to skip straight from here to the section titled **A simple demo** and then come back to examine the example program files in detail, and also find out more about eLua's multi-processing support and ALOHA.

If the demo program has a *serverbg.lua* server, it can be used in place of the normal server when executing the eLua example program, since it implements the same services in addition to executing a background task (which simply flashes a LED on the propeller). For example:

**elua client serverbg**

## Multi-Processing with eLua

This section describes using eLua to execute programs that are not specifically written for eLua. The examples used are the example programs described in the document **Lua on the Propeller 2 with Catalina**, and should already be on the Catalyst SD card (if not, copy the files *ex\*.lua* from *demos/catalyst/lua-5.4.8/test*). However, these examples can be executed as client programs by eLua and can also be executed in conjunction with an eLua server that performs additional functions.

eLua does not *replace* Lua's existing multi-processing capabilities, it *extends* them.

To execute any of these examples, just specify them as eLua client programs, and **null** as the eLua server program<sup>5</sup>. For example:

**elua ex4 null**

or

---

<sup>5</sup> The file *null.lua* is included in all the demo programs. It can be used as either a client or server that does nothing except provide some dummy tables and functions that must be present in all eLua programs..

```
luac -o ex6.lux ex6.lua  
eluax ex6 null
```

These client programs do not themselves use the client/server capabilities of eLua, but because the server processing code will be loaded and executed even if it is not used, some of the demos need minor tweaks due to the reduced Hub RAM and cogs available to eLua clients (and they must all be executed as eLua clients, because eLua servers do not usually support multi-processing capabilities).

Here are the tweaks required:

- ex1.lua* must be executed by **eluax** (not **eluax**) since it requires the Lua parser.
- ex8.lua* can execute at most 5 Lua threads (not 10).
- ex9.lua* can use at most 2 factories (not 4), and can recycle at most 2 workers (not 4).
- ex12.lua* can use at most 2 factories (not 4) and 2 workers (not 4).

Also, note that these examples usually end with a **Press ENTER to terminate** message - but this will only terminate the client and not the server because these programs are not aware there is a server that needs to be terminated. Instead, reset the Propeller to return to the Catalyst prompt.

## Serial ALOHA

ALOHA extends eLua to allow programs to execute clients and servers on different propellers. No changes are required to eLua programs to make use of this capability.

The best way to get started with ALOHA is with an example. In this section, we will walk through a simple - but fully functional - eLua/ALOHA demo<sup>6</sup>. This demo can be executed as an eLua program on one propeller, or as an ALOHA program on two propellers. If it is only going to be executed on a single propeller, the sections below titled **Preparing multiple propellers** and **Building eLua with ALOHA** can be skipped.

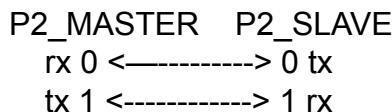
### Preparing multiple Propellers

ALOHA is intended to be used on multiple propellers. The terms 'master' and 'slave' are used for these rather than 'client' and 'server' because when running **eLua**, each propeller runs both a local client and a local server, so we need to distinguish between the client and server running on the master and the clients and servers running on each of the slaves (there can be more than one).

Specifically to support this type of master/slave multi-propeller programming, two new platform configuration files have been added to the *target/p2* folder, and the *target/p2/platforms.inc* file - they are called *P2MASTER.inc* and *P2SLAVE.inc*, which will be used when the Catalina symbols **P2\_MASTER** and **P2\_SLAVE** (respectively) are defined.

The propellers must be connected using two pins for serial communications - the default is to have a dedicated serial connection between the master and each slave, although "multi-dropping" - where a single serial connection is used to connect the master with all the slaves is also possible, provided each slave implements a different set of services.

The 2 or 8 port serial plugin can be used, and the **SIMPLE** HMI option can be used by each propeller (the **TTY** HMI option can be used if the 8 port serial plugin is used, but not if the 2 port serial plugin is used). For this demo we will use the 2 port serial plugin on both the master and slave propellers, with the propellers connected as shown below:



So the file *target/p2/P2MASTER.inc* should specify:

```

' 2 Port Serial constants
' =====
_RX1_PIN    = 0
_TX1_PIN    = 1

```

and the file *target/p2/P2SLAVE.inc* should specify:

```

' 2 Port Serial constants
' =====
_RX1_PIN    = 1

```


---

<sup>6</sup> A more complete version of the demo, with a few more “bells and whistles” is included in **Appendix A**.

```
_TX1_PIN = 0
```

The *P2MASTER.inc* and *P2SLAVE.inc* files included are copies of *P2\_EDGE.inc*, modified as above. If you have a different propeller platform, modify them accordingly.

A picture is probably worth a thousand words at this point, so here is one:



Because the demo program has to be built to execute on two or more different propeller platforms (or two or more of the same platform but perhaps with different options and/or configurations) the easiest way to build them is using Catapult, after setting

**CATALINA\_DEFINE** to the appropriate platform. A **build\_aloha** script to do this is provided (this is described further below).

## Building eLua with Serial ALOHA support

To execute the demo on one Propeller, we do not need to do anything further - we can just execute it with **eluax**. Skip to the next section.

To execute the demo on multiple propellers, we must build eLua for both the master and the slave platforms, and build both the master and slave versions to include ALOHA support.

As with plain eLua, the ALOHA variants differ in the memory models and options. Here is a brief description of the ALOHA variants provided:

<b>aluax</b>	a CMM Lua client with Lua multi-processing support enabled, and an XMM LARGE Lua server with an 8K cache, both <i>with</i> a Lua parser.
<b>aluafx</b>	a CMM Lua client with Lua multi-processing support enabled, and an XMM LARGE Lua server with a 64K cache, both <i>without</i> a Lua parser.
<b>aluafx</b>	an NMM Lua client and an XMM LARGE Lua server with an 8K cache, both without a Lua parser or multi-processing support. Because the client is a NATIVE program it will execute faster than the other eLua variants - but there is less Hub RAM available, so only smaller Lua clients (or servers, depending on how the Hub RAM is allocated) can be executed.

As previously described, building **aluax** (and **aluafx**) can be most easily done using Catapult, after defining **CATALINA\_DEFINE** to specify the appropriate platform, as follows:

```
set CATALINA_DEFINE=P2_SLAVE
catapult alua.c
catapult aluax.c
catapult aluafx.c
```

or

```
set CATALINA_DEFINE=P2_MASTER
catapult alua.c
catapult aluax.c
catapult aluafx.c
```

The **build\_aloha** script in the folder *demos/elua/aloha* does exactly this, but it also renames the resulting P2\_MASTER binaries as *master.bin*, *masterx.bin* and *masterfx.bin*, and the P2\_SLAVE binaries as *slave.bin*, *slavex.bin* and *slavefx.bin* - so we will end up the following six binaries:

<i>master.bin</i>	aluax compiled for P2_MASTER
<i>masterx.bin</i>	aluafx compiled for P2_MASTER
<i>masterfx.bin</i>	aluafx compiled for P2_MASTER

<b>slave.bin</b>	<b>alua</b> compiled for P2_SLAVE
<b>slavex.bin</b>	<b>aluax</b> compiled for P2_SLAVE
<b>slavefx.bin</b>	<b>aluafx</b> compiled for P2_SLAVE

Then **master** (or **masterx** or **masterfx**) must be executed on the P2\_MASTER propeller, and **slave** (or **slavex** or **slavefx**) must be executed on the P2\_SLAVE propeller. The **master** and **slave** names are arbitrary - in fact all the binaries are simply eLua, but built for a specific platform, and including a custom Lua dispatcher and the ALOHA protocol (both described later in this document).

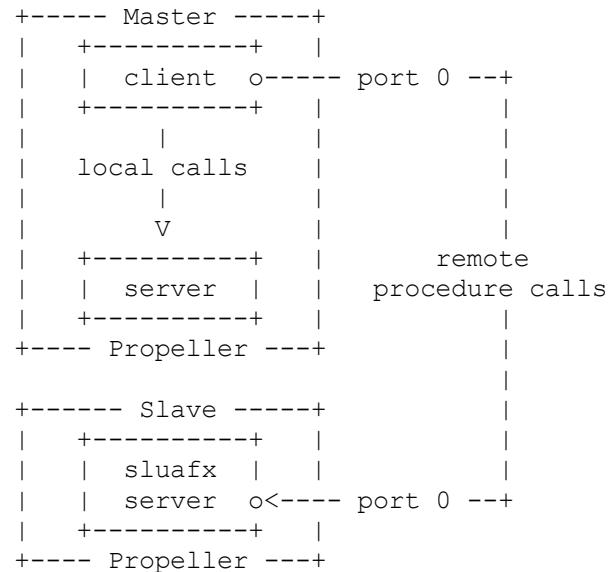
The *alua.c*, *aluax.c* and *aluafx.c* programs are configured to use the 2 port serial plugins by default, but the 8 port serial plugin can be used if more than 3 propellers are to be linked in a star configuration (with one master connected to up to eight slaves - see the diagram above).

To use the 8 port serial plugin, simply specify **-Iserial8** in place of **-Iserial2** before compiling *elua.c*, *eluax.c* or *eluafx.c* (i.e. modify the Catapult pragmas accordingly). Then modify the 8 port serial configuration parameters in *P2MASTER.inc* and *P2SLAVE.inc* platform configuration files (adding more files if there are different slave configurations - e.g. *P2SLAVE0.inc* .. *P2SLAVE7.inc*). It is possible to use a mix of 2 port and 8 port plugins.

There are also some special eLua programs that apply only to Serial ALOHA. These are programs that have only an eLua server, but no eLua client (and so *must* be executed in conjunction with an eLua Serial ALOHA client executing on another Propeller):

- sluafx** an NMM Lua Server with no Lua client, no Lua parser and no multi-processing support. Because the server is a NATIVE program it will execute faster than the other eLua variants, but only support smaller Lua servers.
- sluacx** a CMM Lua Server with no Lua client, no Lua parser and no multi-processing support. Because the server has no client support, it should support larger Lua servers.
- sluacix** a CMM Lua Server with no Lua client, no Lua parser and no multi-processing or floating point support. Because the server has no client support, it should support larger Lua servers.

Another picture might clarify this. In the following picture, the Master Propeller could be executing **elua**, **eluax** or **eluafx**, but there is no client on the Slave Propeller because it is executing **sluafx** (but there may be other slaves running other eLua variants):



The fastest possible combination is when the Master is running **aluafx** and the Slaves are running **sluafx**. In this case the client and the slaves are both executing as NATIVE programs from Hub RAM. **Appendix B** contains an example of using **sluafx**.

## A simple demo

This section describes loading and executing a very simple ALOHA program in detail. The demo can be found in the folder *demos/elua/simple*.

It must be executed on two propellers connected via a serial connection - a **P2\_MASTER** and a **P2\_SLAVE**.

On the P2\_MASTER we will execute *master.bin* and on the P2\_SLAVE we will execute the demo using *slave.bin*

### Loading the simple demo

Copy the appropriate eLua binaries (i.e. *elua.bin*, *master.bin* and *slave.bin*) to the appropriate Catalyst SD cards (for two propellers, one SD card is needed for the P2\_MASTER propeller and one for each P2\_SLAVE propeller).

To each SD card also copy the contents of the *demos/elua/simple* folder, which will include:

<i>binser.lua</i>	the <b>binser</b> module
<i>common.lua</i>	common definitions for the clients and servers
<i>client.lua</i>	an example eLua client
<i>server.lua</i>	an example eLua server
<i>remote.lua</i>	an ALOHA proxy client/server
<i>null.lua</i>	a null client/server
<i>rebuild</i>	a Catalyst script to re-compile all the Lua files

There may also be compiled (i.e. *.lux*) versions of the Lua files, and a *README.TXT* file. Note that these files have the same names as the ones in the eLua example - those files will be overwritten if they already are on the SD card.

Note that if there are no *.lux* versions, or you change any of the *.lua* versions, you should compile the *.lua* files to generate the corresponding *.lux* versions, at least for the common and binser Lua modules. For example:

```
luac -o common.lux common.lua
luac -o binser.lux binser.lua
```

or

```
clua common
clua binser
```

A catalyst script to recompile all the files is provided with each demo called **rebuild**. To execute it, use the following Catalyst command:

```
exec rebuild
```

## Executing the simple demo

To execute the demo on a single propeller, just execute **elua** (or **master** or **slave**) with no parameters.

They will all execute *client.lua* and *server.lua* as the eLua client and server, and the client calls functions in the local server - i.e. in the same propeller. This will not use ALOHA even if it is available. To instead execute the demo on *multiple* propellers, do the following:

First<sup>7</sup>, on the P2\_SLAVE, execute **slave** but specify **remote** as the *client*.

**slave remote server**

Then, on the P2\_MASTER, execute **master** but specify **remote** as the *server*.

**master client remote**

The client on the master propeller will now use ALOHA to call the server on the slave propeller.

## The simple demo files

This section contains a walk-through of all the significant lines in all the files in the *demos/elua/simple* folder. Each file is only a few lines long - to see them all on a single page, see **Putting it all together**, below.

### **client.lua**

The file *client.lua* contains only two lines of significance.

The first loads definitions required by all clients and all servers:

```
dofile 'common.lux'
```

This line simply executes *common.lux*, which is the compiled version of *common.lua* (do not jump ahead just yet to read the section on *common.lua* - read the sections in the order they are included in this document - they are given in this order for a reason).

This line could actually specify *common.lua* - however, it is generally recommended to use the *compiled* version of the common module (i.e. *common.lux*) rather than the *text* version (i.e. *common.lua*) because the former will work even if *client.lua* is subsequently compiled (e.g. to *client.lux*) and then executed by a version of eLua that does not include the Lua parser, whereas the latter will not (it will report an error). However, using the text version is useful while developing programs, as it does not need to be recompiled after each amendment.

The other significant line is the one that shows how the client calls a server function. In the following line, a perfectly ordinary Lua function called **invoke** is being called, and the result of the call is then printed:

```
print(invoker(function(x) return x*x end, 2.5))
```

---

<sup>7</sup> In this simple version the commands *must* be executed in this order - the slave has to be ready *before* the master starts or it will time out. This is not true in the more complete version given in Appendix A.

The **invoke** function may be familiar. It is essentially the same as the one in the non-ALOHA eLua demo. It accepts two parameters:

- a function that accepts a value; and
- a value to be passed to that function.

If this line is a little cryptic, it is very likely because in Lua, functions are *first class types*, and can be specified anywhere any other data type can be specified. Most languages do not support this, and treat functions as second class citizens that must be specially defined - generally *before* they can be used anywhere.

Lua can do that too, of course. The line above could equally well be written as shown below, which may make it clearer that a function **f** that takes a parameter called **x** and returns a result (**x\*x** in this case) is being passed to **invoke**, along with a value for **x**:

```
function f(x)
    return x*x
end
x = 2.5
result = invoke(f, x)
```

Neither the names **f** nor **x** are significant here. Consider the line below - neither **f** nor **x** appear, but this line is also perfectly acceptable. Can you predict what it will do if it is executed as part of *client.lua*?

```
invoke(function(str) print(str) end, "ALOHA"))
```

We can include this line in *client.lua* to find out.

### **server.lua**

The file **server.lua** is a little more complex. But it also contains only a few lines of significance.

Again, the first one loads definitions required by all clients and all servers:

```
dofile 'common.lua'
```

The next significant line defines serial port usage - it defines a table of functions that will be called via the specified serial port.

One such entry is required for each ALOHA port that might call this server, called **port\_0\_index** .. **port\_7\_index**. An empty table indicates the port is to be *monitored* for inward function calls, but no remote calls are made *out* of this port. In our example program, we only want to monitor port 0 for incoming remote calls, so we need only one line:

```
port_0_index = { }
```

The next significant lines are those that define the **invoke** function, which is the function that clients will call remotely:

```
function invoke(serial)
    f, x = bs.deserializeN(serial, 2);
    output = f(x);
    return bs.serialize(output)
```

```
end
```

This is not the normal Lua function definition we might have expected for **invoke**. First, it accepts only one parameter called **serial** (rather than the two parameters **f** and **x** we might have expected). It calls **deserializeN** on that parameter to get the actual parameters **f** and **x**, then it invokes **f** on **x** (which is what we would have expected), and then it calls **serialize** on the output before it returns it.

This (apart from the line which calls the actual function) is how *all* services must be written. Although the service appears to accept just one serialized argument and return one serialized value, those serialized values can actually encode multiple *actual* values.

When the function **invoke** is called by the client, the client actually calls the *proxy* function (defined in **common.lua**) which passes a *serialized* version of all its parameters (which in this case are the 2 parameters **f** and **x**). The function **deserializeN** *deserializes* the specified number of parameters from its input (again, in this case the 2 parameters **f** and **x**). The next line simply invokes the function **f** on parameter **x** and generates output, and the last line calls **serialize** to turn the result into a form that can be passed back to the caller.

This will become clearer when we look at the proxy function for **invoke** defined in **common.lua**.

The important point here is that the function arguments and function return values are always passed between eLua clients and servers - whether those clients are on the same propeller or on different propellers - in a *serialized* binary format. As we have seen, these arguments can be any simple Lua data type, including functions and tables. They are *serialized* before transmission, and must be *unserialized* on reception. This is accomplished using functions provided by the package **binser**, which is a general purpose Lua *serializer* and *deserializer*.

A detailed discussion of **binser** is beyond the scope of this document, but a very quick summary of all we need to know is given in the section titled **The binser module** - and the example above illustrates this is quite easy.

The last significant lines are those that define what the server should do when it is not servicing remote calls. This is a special function called **background**. In this case, we don't need the server to do anything except respond to remote calls to **invoke**, so it is just a null procedure.

```
function background()
end
```

A more detailed discussion of background tasks is given in the section titled **Background tasks**.

### **remote.lua**

This file is only relevant when using ALOHA. In this example program, the file **remote.lua** is almost completely trivial, and contains only a few lines of significance.

Again, the first one loads definitions required by all clients and all servers:

```
dofile 'common.lua'
```

The next significant lines are those that define the port that will be monitored for service calls, and the calls that are expected on that port. In this case, we expect only one call to be made, and they must be made via port 0, to our **invoke** service:

```
port_0_index = {
    [INVOKE_SVC] = "invoke"
}
```

Contrast *this* instance of **port\_0\_index** with the previous instance, described in **server.lua**. That instance was for the *slave* server, which *accepts* calls on port 0 but doesn't make any, and this instance is for the *master* server, which actually makes such calls.

The last significant lines are those that define what the remote server should do when it is not servicing calls. Again, there is a function called **background**. Typically, it will be identical to the one in **server.lua**:

```
function background()
end
```

### **common.lua**

The **common.lua** file is where most of the 'magic' happens.

First, of all, there is some necessary housekeeping:

```
svc = require 'service'
bs = dofile 'binser.lua'
```

The **service** module is required to allow eLua programs to interact with the Catalina Registry. Very briefly, the Registry provides the necessary mechanisms to allow communication between two Lua programs, whether they are running on different cogs on the *same* propeller, or on *different* propellers.

The **binser** module is used by all eLua clients and servers to serialize and deserialize function arguments and return values. See the section titled **The binser module** for more details.

The next significant lines are where we define a Catalina service id for each service we intend to provide (in this case there is only one - for the **invoke** service).

```
INVOKE_SVC = 81
```

Why **81**? Well, all service requests go via the Registry, even if the service is to be provided by another propeller, and Catalina supports service ids in the range 0 .. 255. But services ids 1 .. 80 are reserved for Catalina's own use. So eLua service ids must start from 81.

The next significant lines define an index of services the server will provide, as a Lua table called **service\_index**. In this case, we need only one entry, again for the **invoke** service:

```
service_index = {
    [INVOKE_SVC] = "invoke"
}
```

The next significant lines define a table that describes the WiFi network, and the services that will be called via WiFi. This is not used for *serial* ALOHA services - it is only used by *WiFi* ALOHA services, and will be described in the section on WiFi ALOHA:

```
rpc_network = {
    ["SSID"]      = "MyNetwork";
    ["PASSPHRASE"] = "TellMeASecret";
    [INVOKE_SVC]   = "xxx.xxx.xxx.xxx";
}
```

Finally, we define a *proxy* service for each of the functions we want the client to be able to call. Again, in this case we need only one, for the **invoke** service:

```
function invoke(f, x)
    return bs.deserializeN(svc.serial(INVOKE_SVC, bs.serialize(f, x), 500), 1)
end
```

Note the use of **binser** again here. This instance of **invoke** is a proxy for the actual **invoke** function in **server.lua** (which is where the real function actually executes). This proxy function is what makes the whole thing work - it allows clients to call the **invoke** function exactly as they normally would, without having to know where that function is executed. It may be executed on the local master server, or it may be executed on a remote slave server.

The parameter 500 (to **svc.serial**) specifies how large a buffer may be required to be allocated internally to accommodate the serialize/deserialize process. An error will be reported if the buffer is not large enough.

### **null.lua**

The file **null.lua** is included to provide a null client or server when required. It does nothing, and contains only lines that prevent eLua from complaining about missing tables or functions it expects to find in all clients and servers.

## Putting it all together

Every significant line in the simple demo has been explained in the previous sections. But it may be easier to understand when it is seen all together - so here it is again, in full:

### *common.lua:*

```

svc = require 'service'
bs = dofile 'binser.lux'

INVOKESVC = 81

service_index = {
    [INVOKESVC] = "invoke"
}

rpc_network = {
    ["SSID"]      = "MyNetwork";
    ["PASSPHRASE"] = "TellMeASecret";
    [INVOKESVC]   = "xxx.xxx.xxx.xxx";
}

function invoke(f, x)
    return bs.deserialize(svc.serial(INVOKESVC, bs.serialize(f, x), 500), 1)
end

```

### *client.lua:*

```

dofile 'common.lux'

invoke(function(str) print(str) end, "ALOHA")
print(invoke(function(x) return x*x end, 2.5))

```

### *server.lua:*

```

dofile 'common.lux'

port_0_index = { }

function invoke(serial)
    f, x = bs.deserialize(serial, 2)
    output = f(x)
    return bs.serialize(output)
end

function background()
end

```

### *remote.lua:*

```

dofile 'common.lux'

port_0_index = {
    [INVOKESVC] = "invoke"
}

function background()
end

```

## WiFi ALOHA

For Propellers equipped with a Parallax WiFi module, eLua also allows clients and servers on different propellers to interact using a Remote Procedure Call (RPC) protocol over WiFi. Also, it is possible to use ALOHA serial links for some services and WiFi for others. The next section demonstrates such a hybrid program.

Note that WiFi ALOHA RPC calls are currently limited to a total size of 512 bytes. This limit applies to the sum of all the arguments after they have been serialized by the **binser** module, so it is difficult to be precise about any particular instance, but (for example) if an RPC call took a single string parameter, the string would have to be less than about 508 bytes (to allow for the **binser** overhead).

## Preparing multiple Propellers

Similar to Serial ALOHA, WiFi ALOHA is intended to be used on multiple propellers, which must all be equipped with Parallax WiFi modules. As with Serial ALOHA, there are several platform configuration files provided that support WiFi ALOHA.

As with Serial ALOHA, the terms 'master' and 'slave' are used for these propellers rather than 'client' and 'server'. The 'master' *makes* RPC calls, and the 'slaves' *service* them.

First, for programs that use *only* WiFi, there is a common platform configuration file in the *target/p2* folder called *P2\_WIFI.inc*, which will be used when the Catalina symbol **P2\_WIFI** is defined.

Specifically to support the *hybrid* model of master/slave multi-propeller programming, there are two additional platform configuration files in the *target/p2* folder - they are called *P2WIFI\_M.inc* and *P2WIFI\_S.inc*, which will be used when the Catalina symbols **P2\_WIFI\_MASTER** and **P2\_WIFI\_SLAVE** (respectively) are defined.

For programs that use only WiFi, the Propellers need no physical connection - they just need to be on the same WiFi network. While there can be only one master propeller, there can be any number of slave Propellers in the WiFi network - unlike ALOHA, which is limited to eight slaves.

For hybrid programs that also use serial ALOHA services, there can be up to six slave propellers that use ALOHA serial connections, which must be connected to the master using two pins for serial communications, as described in the ALOHA section.

For programs that use only WiFi, either the 2 or 8 port serial plugin can be used, and the **SIMPLE** HMI option can be used by each propeller (the **TTY** HMI option can be used if the 8 port serial plugin is used, but not if the 2 port serial plugin is used).

For hybrid programs, the 8 port serial plugin must be used, and the serial connections (similar to those used in ALOHA programs) should use serial ports 2 .. 7.

The *P2\_WIFI.inc* and *P2WIFI\_M.inc* and *P2WIFI\_S.inc* files included are copies of *P2\_EDGE.inc*. If you have a different propeller platform, modify them accordingly.

## Building eLua with WiFi ALOHA support

To execute the demo on multiple propellers, we must build eLua for both the master and the slave platforms, and build both the master and slave versions to include WiFi support. This is accomplished by linking with the WiFi library (i.e. `-lwifi`) and one of the serial plugin libraries (i.e. `-lserial2` or `-lserial8`).

As with the Serial variants, the WiFi variants differ in the memory models and options. Here is a brief description of the WiFi (aka RPC) variants provided:

<b>rlua</b>	a CMM Lua client with Lua multi-processing support enabled, and an XMM LARGE Lua server with an 8K cache, both <i>with</i> a Lua parser. Uses the 8 port serial plugin and supports both serial and WiFi ALOHA services. Both the client and the server support the WiFi functions.
<b>rluax</b>	a CMM Lua client with Lua multi-processing support enabled, and an XMM LARGE Lua server with an 8K cache, both <i>without</i> a Lua parser. Uses the 8 port serial plugin and supports both serial and WiFi ALOHA services. Both the client and the server support the WiFi functions.
<b>rlua2</b>	a CMM Lua client with Lua multi-processing support enabled, and an XMM LARGE Lua server with an 8K cache, both <i>with</i> a Lua parser. Uses the 2 port serial plugin and does <b>NOT</b> support serial ALOHA services - only WiFi ALOHA services, and only the server supports the WiFi functions <sup>8</sup> . This allows more Hub RAM for Lua programs.
<b>rlua2x</b>	a CMM Lua client with Lua multi-processing support enabled, and an XMM LARGE Lua server with an 8K cache, both <i>without</i> a Lua parser. Uses the 2 port serial plugin and does <b>NOT</b> support serial ALOHA services - only WiFi ALOHA services, and only the server supports the WiFi functions <sup>9</sup> . This allows more Hub RAM for Lua programs.
<b>rluafx</b>	an NMM Lua client with Lua multi-processing support enabled, and an XMM LARGE Lua server with an 8K cache, both <i>without</i> a Lua parser. Uses the 2 port serial plugin and does <b>NOT</b> support serial ALOHA services - only WiFi ALOHA services, and only the server supports the WiFi functions <sup>10</sup> .

As previously described, building `rlua` (and `rluax`, `rlua2`, `rlua2x` and `rluafx`) can be most easily done using Catapult, after defining `CATALINA_DEFINE` to specify the appropriate platform, as follows:

---

<sup>8</sup> However, the client can use the WiFi functions provided by the server if suitable proxy services are defined - see the HTTP example in `demos/elua/http`.

<sup>9</sup> Ditto.

<sup>10</sup> Ditto.

```
set CATALINA_DEFINE=P2_WIFI
catapult rlua.c
catapult rluax.c
catapult rlua2.c
catapult rlua2x.c
catapult rluafx.c
```

or

```
set CATALINA_DEFINE=P2_WIFI_SLAVE
catapult rlua.c
catapult rluax.c
... etc ...
```

or

```
set CATALINA_DEFINE=P2_WIFI_MASTER
catapult rlua.c
catapult rluax.c
... etc ...
```

The `build_rpc` script in the folder `demos/elua/aloha` does exactly this, but it also renames the resulting P2\_MASTER binaries as `rmaster.bin` and `rmasterx.bin` and the P2\_SLAVE binaries as `rslave.bin` and `rslavex.bin` - to end up the following binaries:

<code>rlua.bin</code>	<code>rlua</code> compiled for P2_WIFI
<code>rluax.bin</code>	<code>rluax</code> compiled for P2_WIFI
<code>rluafx.bin</code>	<code>rluafx</code> compiled for P2_WIFI
<code>rlua2.bin</code>	<code>rlua2</code> compiled for P2_WIFI
<code>rlua2x.bin</code>	<code>rlua2x</code> compiled for P2_WIFI
<code>rmaster.bin</code>	<code>rlua</code> compiled for P2_WIFI_MASTER
<code>rmasterx.bin</code>	<code>rluax</code> compiled for P2_WIFI_MASTER
<code>rslave.bin</code>	<code>rlua</code> compiled for P2_WIFI_SLAVE
<code>rslavex.bin</code>	<code>rluax</code> compiled for P2_WIFI_SLAVE

The `rlua` (or `rluax` or `rluafx` or `rlua2` or `rlua2x`) can be executed on any P2\_WIFI propeller. The `rmaster` (or `rmasterx`) must be executed on the P2\_WIFI\_MASTER propeller, and `rslave` (or `rslavex`) must be executed on the P2\_WIFI\_SLAVE propeller.

The names are arbitrary - in fact all the binaries are simply eLua, but built for a specific platform, and including a custom Lua dispatcher with ALOHA serial and WiFi RPC support (except for `rlua2`, `rlua2x` and `rluafx`, which explicitly disable the ALOHA serial protocol to save Hub RAM).

Just like Serial ALOHA, there is also one special eLua program that applies only to WiFi ALOHA. This is a program that has only an eLua server, but no eLua client (and so *must* be executed in conjunction with an eLua WiFi ALOHA client executing on another Propeller):

**sluarcx** a CMM Lua Server with no Lua client, no Lua parser and no multi-processing support. Because the server has no client support, it should support larger Lua servers.

## A hybrid demo

This section describes loading and executing a hybrid ALOHA program which uses both ALOHA serial and WiFi services. The demo can be found in the folder *demos/elua/hybrid*.

It must be executed on two propellers, each with a WiFi module, and also connected via a serial connection - a **P2\_WIFI\_MASTER** and a **P2\_WIFI\_SLAVE**.

Of course, there is no real need for a serial connection if both Propellers have WiFi adaptors - this is done only for this hybrid demo, to show that this can be supported. All the other example eLua programs provided can also be executed using two propellers connected only via WiFi.

On the P2\_WIFI\_MASTER we will execute *rmaster.bin* and on the P2\_WIFI\_SLAVE we will execute the demo using *rslave.bin*

### Loading the hybrid demo

Copy the appropriate eLua binaries (i.e. *rmaster.bin* and *rslave.bin*) to the appropriate Catalyst SD cards (one SD card is needed for the P2\_MASTER propeller and one for each P2\_SLAVE propeller).

The P2\_MASTER and P2\_SLAVE propellers must have their WiFi modules installed on the pin groups specified in the *P2\_WIFI\_M.inc* and *P2\_WIFI\_S.inc* platform configuration files, and the two propellers must also have a serial connection connecting the pins defined for the third serial port (by default these are pins 0 and 1).

The WiFi modules on each Propeller may need to have their firmware updated - see the *README.TXT* file in the *demos/wifi* folder for more details.

To each SD card also copy the contents of the *demos/elua/hybrid* folder, which will include:

<i>binser.lua</i>	the <b>binser</b> module
<i>common.lua</i>	common definitions for the clients and servers
<i>client.lua</i>	an example eLua client
<i>server.lua</i>	an example eLua server
<i>remote.lua</i>	an ALOHA proxy client/server
<i>null.lua</i>	a null client/server
<i>rebuild</i>	a Catalyst script to re-compile all the Lua files

Note that if there are no *.lux* versions, or you change any of the *.lua* versions, you should compile the *.lua* files to generate the corresponding *.lux* versions, at least for the common and binser Lua modules. For example:

```
luac -o common.lux common.lua
luac -o binser.lux binser.lua
```

or

**clua common**  
**clua binser**

A catalyst script to recompile all the files is provided with each demo called **rebuild**. To execute it, use the following Catalyst command:

**exec rebuild**

## Executing the hybrid demo

The hybrid demo must be executed on both the **P2\_WIFI\_MASTER** and **P2\_WIFI\_SLAVE** propellers:

First<sup>11</sup>, on the **P2\_WIFI\_SLAVE**, execute **rslave** but specify **remote** as the *client*:

**rslave remote server**

Then, on the **P2\_WIFI\_MASTER**, execute **rmaster** but specify **remote** as the *server*:

**rmaster client remote**

The client on the master propeller is using WiFi to call the server on the slave propeller for all services except **Invert**, which will use the serial ALOHA connection.

Both programs will print WiFi information on startup, which can be used to determine the IP addresses of the Propellers once they have joined the WiFi network. For example, you might see something like:

```
WiFi SSID = 'MyNetwork'
WiFi module name = 'wx-XXXXXX'
WiFi JOIN 'MyNetwork'
WiFi IP address = '192.168.1.115'
```

The **SSID** will depend on what is specified in the **rpc\_network** table in *common.lua*, and the program will only JOIN successfully and display an IP address if the **PASSPHRASE** is also correct.

Also, the program will not JOIN a network if the **SSID** is either an empty string or the same as its own WiFi module name - in that case, it assumes it is the one offering its **SSID** as an Access Point, and all the other propellers in the network should JOIN this Access Point. Only the master propeller should be configured this way.

The **IP address** displayed will depend on your local DHCP server, and the program should be run once to obtain the IP address, which can then be used to specify the IP address of the server offering each RPC service in the **rpc\_network** table in *common.lua* (and then **exec rebuild** the program).

The output of this WiFi information on startup can be disabled by setting **WIFI\_INFO** to 0 in *dsptch\_l.c* and then recompiling the eLua programs.

---

<sup>11</sup> In this demo the commands *must* be executed in this order - the slave has to be ready *before* the master starts or it will time out.

## A HTTP demo

This section describes loading and executing a WiFi/RPC program where the master propeller serves a web page, and retrieves data from a slave propeller using WiFi/RPCs to display information on that web page. The demo can be found in the folder *demos/elua/http*.

It is intended to demonstrate that a propeller can both *respond* to HTTP requests (i.e. to display the web page), and also *make* HTTP requests to fetch data to display from other propellers<sup>12</sup>.

It must be executed on two propellers, each with a WiFi module.

We will execute the demo using *rLuaX.bin* on both the master and slave propellers.

### Loading the HTTP demo

Copy the appropriate eLua binaries (i.e. *rLuaX.bin*) to the appropriate Catalyst SD cards (one SD card is needed for the master propeller and one for the slave propeller).

The **master** and **slave** propellers must have their WiFi modules installed on the pin groups specified in the *P2\_WIFI.inc* platform configuration file.

The WiFi modules on each Propeller may need to have their firmware updated - see the *README.TXT* file in the *demos/wifi* folder for more details.

To each SD card also copy the contents of the *demos/elua/http* folder, which will include:

<i>binser.lua</i>	the <b>binser</b> module
<i>common.lua</i>	common definitions for the clients and servers
<i>client.lua</i>	the eLua client
<i>server.lua</i>	the eLua server
<i>serverbg.lua</i>	the eLua server with a background task
<i>remote.lua</i>	an ALOHA proxy client/server
<i>null.lua</i>	a null client/server
<i>rebuild</i>	a Catalyst script to re-compile all the Lua files

Note that if there are no *.lux* versions, or you change any of the *.lua* versions, you should compile the *.lua* files to generate the corresponding *.lux* versions. A catalyst script to recompile all the files is provided with each demo called **rebuild**. To execute it, use the following Catalyst command:

**exec rebuild**

---

<sup>12</sup> The “Same Origin Policy” of modern browsers would normally prevent such a simple web page from displaying information from another “origin” (i.e. a different propeller) on the same web page. This program gets around that by retrieving the information from the other propeller using an RPC, which it then sends back to the browser.

## Executing the HTTP demo

The HTTP demo must be executed on both the **master** and **slave** propellers. For example:

On the **slave**, execute **rluax** but specify **remote** as the *client*:

**rluax remote server**

On the **master**, execute **rluax** but specify **remote** as the *server*:

**rluax client remote**

The client on the **master** propeller will use WiFi/RPC to call the server on the **slave** propeller to retrieve data to display on the web page.

Both programs will print WiFi information on startup, which can be used to determine the IP addresses of the Propellers once they have joined the WiFi network. For example, you might see something like:

```
WiFi SSID = 'MyNetwork'
WiFi module name = 'wx-XXXXXX'
WiFi JOIN 'MyNetwork'
WiFi IP address = '192.168.1.115'
```

The **SSID** will depend on what is specified in the **rpc\_network** table in *common.lua*, and the program will only JOIN successfully and display an IP address if the **PASSPHRASE** is also correct.

Also, the program will not JOIN a network if the **SSID** is either an empty string or the same as its own WiFi module name - in that case, it assumes it is the one offering its **SSID** as an Access Point, and all the other propellers in the network should JOIN this Access Point. Only the master propeller should be configured this way.

The **IP address** displayed will depend on your local DHCP server, and the program should be run once to obtain the IP address, which can then be used to specify the IP address of the server offering each RPC service in the **rpc\_network** table in *common.lua* (and then **exec rebuild** the program).

The output of this WiFi information on startup can be disabled by setting **WIFI\_INFO** to 0 in *dsptch\_l.c* and then recompiling the eLua programs.

The master propeller will also print a message like:

**Open a browser to <http://192.168.1.117/prop>**

Do as it says, and then follow the instructions on the displayed web page.

An interesting aspect of this HTTP demo is that although the master propeller **client** contains the web page to serve (which is a simple string in this demo program, but would normally be loaded from a file), it is the master propeller **server** that does all the WiFi processing. Examine the **service\_index** table in *common.lua*. It will contain entries like:

```
service_index = {
    [LISTEN_SVC] = "listen",
```

```
[POLL_SVC]    = "poll",
[SEND_SVC]    = "send",
[CHECK_SVC]   = "check",
... etc ...
[INVOKE_SVC]  = "invoke",
[QUIT_SVC]    = "quit"
}
```

Also, examine the `rpc_network` table in `common.lua`:

```
rpc_network = {
  ["SSID"]        = "MyNetwork";
  ["PASSPHRASE"] = "TellMeASecret";
  [INVOKE_SVC]   = "xxx.xxx.xxx.xxx";
  [QUIT_SVC]     = "xxx.xxx.xxx.xxx";
}
```

What this indicates is that the client uses `listen`, `poll` and `send` WiFi functions implemented in the local server to handle the HTTP requests, but `invoke` is implemented in the remote propeller, and the client uses a WiFi/RPC call to access it.

Implementing complex functions in the server rather than the client is typical of ALOHA programs (whether serial or WiFi). The client is limited in size to what is available in Hub RAM, but the server can have access to many megabytes of data and code space available in XMM RAM.

# Technical Notes

## File name conventions

An eLua program with one client and one server typically consists of three files. They can be given any names, but the following convention is recommended:

<b>common.lua</b>	definitions common to all the clients and servers in the eLua program
<b>client.lua</b>	the master client Lua program.
<b>server.lua</b>	the slave server Lua program.

eLua programs that use ALOHA can span multiple propellers, and at least one more file is typically required:

<b>remote.lua</b>	the master server Lua program, which often may also conveniently be used as the slave client program.
-------------------	---

If there is more than one slave (i.e. a multi-propeller eLua program with more than two propellers), then adding suffixes **\_0 .. \_7** to the server and remote file names is recommended.

Following this convention means that the common file is the only one that needs to know the details of the service ids required to implement the client/server communications, and also provides a place to implement the proxy services that allow the client to be unaware that the services are not implemented locally. It also provides a convenient place to load the common modules (e.g. **binser**) and perform any necessary client or server configuration.

The service ids (a unique one must be allocated to each service) can be any service id other than those reserved for Catalina's own use (see *target/p2/constants/inc*) - i.e. from (**SVC\_RESERVED + 1**) to 255 - ii.e. 81 to 255.

## Lua initialization module

The main job of the Lua initialization module (*linit.c*) is to load the required Lua libraries. It does so based on various symbol definitions, which may be defined explicitly in some cases, or by which Catalina libraries are specified in others. The significant ones for eLua/ALOHA are:

<b>-luax</b>	If the program is compiled using the <b>luax</b> library rather than the <b>lua</b> library (or the C symbol <b>__CATALINA_libluax</b> is defined) then the Lua parser will <i>not</i> be loaded.
<b>-lthreads</b>	If the program is compiled with the <b>threads</b> library (or the C symbol <b>__CATALINA_libthreads</b> is defined) then the Lua threads library will be loaded and the Lua Multi-processing extensions will be enabled.

<b>-Iwifi</b>	If the program is compiled with the <b>wifi</b> library (or the C symbol <b>_CATALINA_libwifi</b> is defined) then the Lua wifi library will be loaded.
<b>-Iserial2</b>	If the program is compiled with the <b>serial2</b> library (or the C symbol <b>_CATALINA_libserial2</b> is defined) then the Lua serial library will be loaded, and 2 serial ports will be available.
<b>-Iserial8</b>	If the program is compiled with the <b>serial8</b> library (or the C symbol <b>_CATALINA_libserial8</b> is defined) then the Lua serial library will be loaded and 8 serial ports will be available.
<b>-C LUA_NO_MATH</b>	If the Catalina symbol <b>LUA_NO_MATH</b> is defined, then the Lua math library will <i>not</i> be loaded.
<b>-C LUA_NO_OPT</b>	If the Catalina symbol <b>LUA_NO_OPT</b> is defined, then the Lua os, utf8 and debug libraries will <i>not</i> be loaded.
<b>-C LUA_SERVICE</b>	If the Catalina symbol <b>LUA_SERVICE</b> is defined, the <b>services</b> module which is used to implement Lua client/servers will be loaded.
<b>-C LUA_LINENOISE</b>	If the Catalina symbol <b>LUA_LINENOISE</b> is defined, the <b>linenoise</b> library will be loaded.

## Hub RAM Allocation

The Catapult 'stack' and 'address' attributes are the mechanism used to allocate Hub RAM between the client and the server. Trial and error must be used to determine the appropriate sizes. The minimum stack required by ANY client is about 80,000 bytes (which is the size used by the **eluas** version of the demo).

Note that since the client always executes entirely from Hub RAM, the client stack size (specified as the 'stack' attribute of the secondary catapult pragma) includes the Lua stack and heap. Once this size is determined, the resulting server address (specified as the 'address' attribute of the primary catapult pragma) determines how much Hub RAM is left for the server. However, because the server is an XMM LARGE program, this Hub RAM is used only for the Lua stack and does not include the heap (which is in XMM RAM), so the Hub RAM requirement of Lua code executed in the server is typically much lower than it would be if the same code was executed by the client. Put in simple terms, if there is a choice then Lua code should be put in the server rather than the client..

Here are the approximate amounts of Hub RAM (in bytes) available to the client and the server in each eLua variant:

name	client	server
<b>=====</b>	<b>=====</b>	<b>=====</b>
<b>eluax</b>	100,000	100,000
<b>eluax</b>	100,000	90,000
<b>elua</b>	80,000	175,000

<b>eluasx</b>	100,000	130,000
<b>eluafx</b>	100,000	50,000

Other Hub RAM trade-offs between client and server are of course possible.

## Cog allocation

In order to maximize the available free cogs for demonstration purposes, the RTC plugin is not loaded by any of the eLua variants. This may make some Lua functions not work as expected - e.g. **os.date()** or **os.clock()**. If these are required, add the options **-C RTC** (in place of **-C CLOCK**) to the appropriate Catapult pragmas. This requires an additional cog.

## The binser module

The **binser**<sup>13</sup> module does not originate with Catalina - the original is available on github<sup>14</sup>. However, Catalina's version has been modified to suit Catalina and this version is in the folder *demos/elua/binser*, along with the original documentation.

The **binser** module is normally loaded by the common module, and it is recommended to use the compiled version (*binser.lux*) because it loads faster and will also work in compiled clients and servers (the text version *binser.lua* will not). For example:

```
bs = dofile 'binser.lux'
```

There are only three **binser** functions that are typically required by an eLua program:

The function **serialize** is used to turn a series of one or more arguments into a serial binary string (note the result is a Lua string and not a C string - it may contain embedded zeroes):

```
serial = bs.serialize(args ...)
```

For example:

```
serial = bs.serialize(x, y, z)
```

The function **deserializeN** is used to turn a serial binary string back into **N** arguments (if there is only one argument, specify 1 for **N**):

```
arg1, ... argN = bs.deserializeN(serial, N)
```

For example:

```
arg = bs.deserializeN(serial, 1) -- one argument
arg1, arg2 = bs.deserializeN(serial, 2) -- two arguments
```

Finally, there is also a **deserialize** function:

```
args = bs.deserialize(serial)
```

However, **deserialize** returns its results as a Lua table instead of a series of Lua data types, which makes it slightly more complex to use. For instance, instead of using **deserializeN**, to use **deserialize** to do the same job would have to be written as:

```
args = bs.deserialize(serial)
args1 = args[1]
...
argsN = args[N]
```

Using **deserializeN** is easier unless the program *intends* to pass Lua tables, in which case it is easier to use **deserialize**. For instance:

```
args_in = {x=1, y=2} -- this is a table with two elements (x and y)
serial = bs.serialize(args_in)
args_out = bs.deserialize(serial)
```

<sup>13</sup> In all the eLua program examples, **binser** is abbreviated to **bs** when loaded.

<sup>14</sup> See <https://github.com/bakpakin/binser>

There is a Lua subtlety here that can catch you unawares. Both the `deserializEN` and `deseserialize` functions actually return more than one value, which is a perfectly acceptable thing to do in Lua. The first values returned are the results of deserializing, and the final value is the resulting position in the string being deserialized - something which is generally of little interest in eLua.

The subtlety is that if you pass the result of deserializing straight to a function that *accepts* a variable number of arguments - such as the `print` function does - you will get *all* the values.

So, for example:

```
result = bs.deserializEN(bs.serialize("abc"),1)
print(result)
```

will output:

```
abc
```

which is what you would expect. Whereas ...

```
print(bs.deserializEN(bs.serialize("abc"),1))
```

which looks like it should do exactly the same thing, will instead output:

```
abc 6
```

which may be a little surprising until you know what is going on!

## The Serial ALOHA protocol

The Serial ALOHA protocol is a simple serial protocol specifically designed to be used for client/server transactions.

The protocol is asymmetric. The server never initiates a transaction. The client initiates every transaction by sending a request packet:

```
FF 02 id sq lo hi b1 .. bn ck
```

Where:

<code>id</code>	is a service id
<code>sq</code>	is a sequence number (can be used to detect repeats)
<code>lo</code> and <code>hi</code>	are the length of data - i.e. ( <code>hi</code> <<8)+ <code>lo</code>
<code>b1 .. bn</code>	are the bytes of data (may contain zeroes)
<code>ck</code>	is the checksum of packet (excluding the initial <code>FF 02</code> )

The server can respond to the request packet by sending a return `FF 02` packet, which indicates success:

```
FF 02 id sq lo hi b1 .. bn ck
```

Or the server can respond by sending a failure response. Some possible failure responses are:

- FF 01** timeout on tx
- FF 03** checksum error on tx
- FF 04** no such id
- FF 05** other error

Notes:

- Any occurrence of **FF** other than in the initial **FF 02** is "byte stuffed" to **FF 00**. This means that **FF 02** can never occur within a message, and so it always signals the start of a packet.
- The **ck** is set so that the sum of all the bytes after the **FF 02** equals zero (modulo 0x100).
- The **sq** is not checked by the protocol - it is just a value. Typically it is simply returned by the server in the response to the request. If the service is not idempotent then it is up to the server to ensure it does not service duplicate requests, and it is up to the client to ensure the response packet it receives has the correct **sq**, incrementing it as required to ensure the the response is not simply a duplicate that has been buffered and/or re-transmitted during the serial processing.

The protocol is implemented in the files *aloha.h* and *aloha.c* in *demos/elua/aloha*. This folder also contains a simple demo/test program implemented in the files *amaster.c* and *aslave.c* - see those files for more details.

There are only two functions required to implement the ALOHA protocol. These are described in the sections below.

## aloha\_tx

Here is the C function prototype of **aloha\_tx**, which is used to transmit an ALOHA packet:

```
void aloha_tx(int port, int id, int sq, int len, char *buf);
```

- port** 0 .. 1 if the 2 port serial plugin is used, 0 .. 7 if the 8 port serial plugin is used.
- id** the id (0 .. 255) of the service to be called.
- sq** a sequence number that should be incremented on each call. This number will be returned in the response packet so that it is possible for the sender to distinguish a response to a request from a buffered or retransmitted response to a previous request.
- len** the length of the message to send.
- buf** a pointer to the binary message to send.

## aloha\_rx

Here is the C prototype of **aloha\_rx**, which is used to receive an ALOHA packet:

```
int aloha_rx(int port, int *id, int *sq, int *len, char *buf, int max, int ms);
```

**port** 0 .. 1 if the 2 port serial plugin is used, 0 .. 7 for the 8 port serial plugin is used.

**id** a pointer to an int used to return the id (0 .. 255) of the service that was requested.

**sq** a pointer to an int used to return the sequence number of the request.

**len** a pointer to an int used to return the size in bytes of the response data.

**buf** a pointer to a buffer to put the response data.

**max** the size of the buffer.

**ms** the timeout in milliseconds to wait for each byte of the response.

The return value is as follows:

- 0 success (success packet received)
- 1 timeout on rx
- 2 packet larger than max
- 3 checksum error on rx
- n failure packet n received (n != 2)

## WiFi/RPC network definition

The WiFi/RPC network is described in a Lua table called **rpc\_network**, which should be common to all servers. It is recommended that it be included in the **common.lua** module, but it can instead be included in both the **server.lua** and **remote.lua** modules.

This table contains the following keys and key types, and the values have the following meaning:

Key Value	Key Type	Value and Meaning
"SSID"	string	The SSID this Propeller is to join. If this propeller is to offer its SSID as the Access Point other Propellers must join, this should be either an empty string, or the same as the WiFi module name.
"PASSPHRASE"	string	The passphrase required to join the WiFi Access Point. For Access Points offered by propellers, this should be an empty string.
SERVICE_ID	integer	The IP address of the client that offers the service with the specified ID. Each of these integer IDs must be present in the <b>service_index</b> table.

Here are the actual values of the service ids, and the **service\_index** and **rpc\_network** tables from the *hybrid* example:

```
-- define a unique service id for each service:
ADD_SVC      = 81
DIVMOD_SVC   = 82
INVERT_SVC   = 83
INVOKE_SVC   = 84
QUIT_SVC     = 89

-- define the services implemented in the server:
service_index = {
    [ADD_SVC]      = "add",
    [DIVMOD_SVC]   = "divmod",
    [INVERT_SVC]   = "invert",
    [INVOKE_SVC]   = "invoke",
    [QUIT_SVC]     = "quit"
}

-- define the WiFi network and the RPC calls supported:
rpc_network = {
    ["SSID"]        = "MyNetwork";
    ["PASSPHRASE"] = "TellMeASecret";

    [ADD_SVC]      = "xxx.xxx.xxx.xxx";
    [DIVMOD_SVC]   = "xxx.xxx.xxx.xxx";
    [INVOKE_SVC]   = "xxx.xxx.xxx.xxx";
    [QUIT_SVC]     = "xxx.xxx.xxx.xxx";
}
```

The **SSID** and **PASSPHRASE** must be replaced with those of the WiFi network that the propellers are to join. If the master propeller is to offer an Access Point that the slave propellers must join, then that propeller should have both the **SSID** and **PASSPHRASE** set to an empty string, and the other propellers should specify the module name of the master propeller as the **SSID** they will join.

The IP addresses (shown in the table as “**xxx.xxx.xxx.xxx**” must be replaced with the actual IP addresses assigned to the master and slave propellers (which can be obtained by executing the programs - they will be printed on startup). If the IP address specified is an empty string, an invalid IP address, or matches the IP address of the propeller, it means this propeller is a slave and will offer the service via HTTP (and listen for HTTP requests). If it does not match, it indicates that this propeller is a master, and will send a HTTP RPC request to the nominated IP address to service the request.

Note in the above example that the “**invert**” service is included in the **service\_index** table but not mentioned in the **rpc\_network** table. This is a deliberate omission, and means that this service is not available via WiFi - it is only available via an ALOHA serial connection. To identify the serial port that will be used to service the request, the **remote.lua** module contains the following information:

```
-- this is used by ALOHA ...
port_2_index = {
    [ADD_SVC]      = "add",
    [DIVMOD_SVC]   = "divmod",
    [INVERT_SVC]   = "invert",
    [INVOKE_SVC]   = "invoke",
}
```

And the **server.lua** module contains the following information:

```
-- this is used by ALOHA ...
port_2_index = { }
```

The **rpc\_network** table takes precedence over the **port\_n\_tables**, this combination means that the “**add**”, “**divmod**”, and “**invoke**” services will use WiFi, but the “**invert**” service will continue to use serial ALOHA.

## The WiFi ALOHA protocol

The WiFi ALOHA protocol is a very simple protocol designed to be used for RPC transactions across HTTP. The server does a HTTP POST request to the path **/rpc/service\_name** at the IP address specified as the client in the **rpc\_network** table, and waits for a HTTP response.

Since HTTP is essentially a text protocol the binser data sent in both directions is first encoded as base64 data in the body of the POST request before transmission, and sent as a REPLY and decoded on reception.

The base64 encoding and decoding is done by functions provided in all the Catalina libraries - see the include file *base64.h* for details.

## The ALOHA dispatcher

eLua normally uses the Lua dispatcher which is in the standard Catalina C library and is called **\_dispatch\_Lua\_bg**. This dispatcher works perfectly well when the client and server are executing in different cogs on the same propeller - the transfer of data between the client and server can occur in Hub RAM.

However, eLua was designed to be extendable across *multiple* propellers - i.e. when the client is executing on one propeller and one or more servers are executing on another. This is where the ALOHA serial and WiFi protocols come in.

To use the ALOHA serial or WiFi protocols, a custom Lua dispatcher is required. This is called **my\_dispatch\_Lua\_bg** and is provided in *dsptch\_l.c* in the folder *demos/elua/aloha*. Using this dispatcher requires that the 2 port serial or 8 port serial plugin is used, and that the serial ALOHA protocol is also included. Support for the WiFi ALOHA protocol is optional, and is determined by whether the program includes the WiFi library (i.e. **-lwifi**).

This custom version of the Lua dispatcher has the following additions to the standard Lua dispatcher:

- It retrieves the list of all Lua services (local or remote) from the **service\_index** table, and the ports any remote services must use from the **port\_0\_index .. port\_n\_index** tables (where **n** = 1 or 7, depending on which serial plugin is in use). All services that might need to be dispatched on this propeller must be listed in the **service\_index** table. An entry that is listed in the **port\_n\_index** tables (and it can be listed only in *one* of them) is assumed to be a remote service available on port **n**. Any entry that is listed in the **service\_index** table but *not* in any of the **port\_n\_index** tables is assumed to be a local service.

- It retrieves the names and IP addresses of any RPC services from the **rpc\_network** table. All the services in the **rpc\_network** table should also be in the **service\_index** table, and are assumed to be WiFi/RPC services, offered as a WiFi/RPC service by the propeller with the specified IP address (see the **WiFi/RPC network definition** section below).
- It monitors the Catalina Registry as usual for service requests. If WiFi support is enabled and it receives a request for a service whose id corresponds to one in the **rpc\_network** table it sends the request as a HTTP request to the specified IP address and waits for a HTTP response, otherwise if the request corresponds to an id in the **port\_n\_index** table it sends the request as an ALOHA request out the appropriate serial port, and waits for a response on the same port. Otherwise, it dispatches the service internally (i.e. to the cog indicated in the Catalina Service Registry)
- It monitors all ports which have a **port\_n\_index** table (even if that table is empty) for the arrival of serial ALOHA messages. If it receives one, it services the request and returns an ALOHA response on the same port.
- If WiFi support is enabled It listens for HTTP requests on port 80, and if it receives one, it services the request and sends a HTTP response.

It is worth noting that ALL services must be included in the **service\_index** table, but only services that must be called using serial ALOHA should be included in the **port\_n\_index** table.

It is also worth noting that while it is highly recommended that all clients and servers use the same **service\_index** table, each client and server may use a different set of **port\_n\_index** tables. When multiple propellers are configured in a simple star configuration (i.e. with the master propeller connected via a dedicated serial link to each slave propeller) then it is recommended that all propellers use the *same* set of tables (which would typically be specified in **common.lua**). An example of this is given in the demo program. More complex configurations are possible, but are beyond the scope of this document.

The dispatcher has the following configurable parameters, which may need to be adjusted for specific applications:

Parameter	Default	Meaning
REMOTE_TIMEOUT	1000	timeout on response after sending request
PORT_TIMEOUT	200	timeout on rx after receiving first byte
REMOTE_MAX	2048	maximum size of message that can be received (note that this must be <b>wifi_DATA_SIZE</b> if WiFi RPC support is enabled).
IP_RETRIES	30	number of times to retry fetching IP address on network JOIN

IP_RETRY_SECS	3	interval (seconds) between retries of fetching IP address
RETRY_INTERVAL	200	interval (milliseconds) between polls for a reply to an RPC request
REPLY_RETRIES	25	number of times to poll for a reply to an RPC request

## Background tasks

The Lua dispatcher executes on each propeller and dispatches calls made via the registry to their appropriate destination (i.e. either to the local slave server, or to a remote slave server). It also continually checks for incoming service calls on the serial ports and dispatches those as well.

But when there are no service calls outstanding, it can also periodically execute a **background** Lua function. It does this between each check of the registry and the nominated serial ports.

The consequence of this is that if the background function adheres to some simple conditions, then it can be used to perform useful tasks. The **background** function must adhere to the following conditions:

1. It cannot accept any arguments or return any values.
2. It must return regularly, so as not to hold up incoming service calls unduly.
3. If ALOHA is in use, it must not execute for long enough on each call to cause an ALOHA timeout - if it executes for a significant portion of the REMOTE\_TIMEOUT specified for the ALOHA protocol, then a remote service call may fail with a timeout error.

The function should be written so that on each call, it picks up where it left off from the previous call, does a little more processing, and then returns. Note that it can store state internally between calls.

Here is a trivial example of a background task that flashes an LED periodically:

```
count = 0
LED = 38 -- suitable for P2_EDGE, change to 56 for P2_EVAL
function background()
    count = count + 1
    if count > 1000 then
        propeller.togglepin(LED)
        count = 0
    end
end
```

The **background** function can be defined in **server.lua** for automatic execution by the ALOHA dispatcher on slave servers, or **remote.lua** for automatic execution on the master server.

The background function can also be executed on slave clients, but this is not done automatically. To do it manually, simply add a loop to call **background()** repeatedly. It is

recommended to add a small delay between calls, or call **propeller.msleep(0)** between calls (which does a *yield* if multi-threading is in use).

For example, add the following code to **remote\_n.lua** for each slave client that should execute the background function:

```
while true do
    background()
    propeller.msleep(0)
end
```

## Appendix A - A more complete demo

This appendix contains a more complete and “user friendly” version of the simple demo. This version is in the folder `demos/elua/complete`. Copy the files in this folder to a Catalyst SD card to execute them. Note that they have the same names as the files in the simple version of the demo, and will overwrite them.

No new concepts or features are introduced, but it includes comments, does more error checking, has more examples of some features, and has the following improvements over the simple version described in the main section of this document:

1. It prints messages when the client and server Lua programs are loaded.
2. It does not require the slave to be executed first. The master will prompt for a keystroke before proceeding.
3. It demonstrates both *local* and *remote* services, rather than just remote ones (the “quit” service is defined in the `service_index` table in `common.lua`, but not listed in the `port_0_index` table in `remote.lua` - instead, it is defined in both `server.lua` and `remote.lua`, which means calls to this service will always go to the server local to the client that calls it, not to a remote server).
4. It demonstrates a service (the “big” service) that accepts and returns a string significantly longer than the size of the buffers used by the serial plugin. Note that to accommodate this, the proxy function for “big” in `common.lua` specifies a larger buffer size (1000 bytes).
5. It demonstrates a means whereby the master can shut down *remote* clients and servers (using the “execute” service) as well as *local* servers (using the “quit” service).

Execute this program the same way as the simpler version of the demo.

On the master propeller:

**master client.lua remote.lua**

On the slave propeller:

**slave remote.lua server.lua**

### **client.lua**

```

print("client ...")

-- load common definitions ...
dofile 'common.lua'

print("... loaded")

function ENTER_to_continue()
    print("\nPress ENTER to continue");
    io.read();
end

-- call some server functions ...

ENTER_to_continue()

```

```

print("calling invoke ...")
invoke(function(str) print(str) end, "ALOHA")

ENTER_to_continue()
print("calling invoke ...")
function f(x)
    return x*x
end
print ("Result = "... invoke(f, 2.5))

ENTER_to_continue()
print("calling big ...")
input = "Now is the time for all good men to come to the aid of the party
... The quick brown fox jumps over the lazy dog ... Four score and ten
years ago ... we will fight them on the beaches ... hey ho, hey ho - it's
off to work we go ... "
print ("Result = "... big(input .. input .. input .. input))

ENTER_to_continue()
print("calling execute (will time out!)...")
execute("")
print("calling quit ...")
quit()

```

## server.lua

```

print("server ...")

-- load common definitions ...
common = dofile('common.lux')

-- this is used by ALOHA for serial services ...
port_0_index = { } -- empty table means monitor the port for serial calls

-- this is used by ALOHA for WiFi services ...
rpc_network = {
    ["SSID"]      = "MyNetwork";      -- set to "" if propeller offers AP
    ["PASSPHRASE"] = "TellMeASecret"; -- set to "" for a propeller AP

    [INVOKE_SVC]  = "xxx.xxx.xxx.xxx";
    [QUIT_SVC]    = "xxx.xxx.xxx.xxx";
    [BIG_SVC]     = "xxx.xxx.xxx.xxx";
    [EXECUTE_SVC] = "xxx.xxx.xxx.xxx";
}

-- this is the function the server executes in the background ...
function background()
    -- the default is to do nothing!
end

-- define local services ...

-- this service returns the output of invoking function f on value x ...
function invoke(serial)
    f, x = bs.deserializeN(serial, 2);
    output = f(x);
    return bs.serialize(output)
end

-- this service demonstrates accepting and returning a big string ...
function big(serial)

```

```

        input = bs.deserializeN(serial, 1);
        output = "bigger ... " .. input
        return bs.serialize(output)
    end

    -- this service executes a Catalyst command, which will also
    -- terminate the slave ...
    function execute(serial)
        command = bs.deserializeN(serial, 1);
        if type(command) == "string" then
            print("Client requested execute '" .. command .. "'")
            propeller.execute(command)
        else
            result = "Invalid execute command";
        end
        return bs.serialize(output)
    end

    -- this service quits the local server ...
    function quit()
        print("Client requested shutdown\n")
        os.exit()
    end

    print("... loaded")

```

## common.lua

```

-- this file must be loaded by both the client and the server. It is the
-- only file required to know about the service ids.
--

-- for example, to load the text version:
--

--     dofile 'common.lua'
--

-- or to load the compiled version:
--

--     dofile 'common.lux'
--

-- using the compiled version is recommended even in text Lua files,
-- so that the files do not have to be modified to be executed using a
-- client/server program that does not load a Lua parser.

svc = require 'service'
bs  = dofile 'binser.lux'

-- define a unique service id for each service:

INVOKE_SVC      = 81
QUIT_SVC        = 82
BIG_SVC         = 83
EXECUTE_SVC     = 84

-- define the services implemented by the server ...
service_index = {
    [INVOKE_SVC]    = "invoke",
    [QUIT_SVC]      = "quit",
    [BIG_SVC]        = "big",
    [EXECUTE_SVC]   = "execute",
}

```

```
-- define proxy calls for the services implemented by the server. The
-- value of 500 in these definitions is the maximum size (in bytes) that
-- is expected when the parameters are serialized:

function invoke(f, x)
    out = bs.deserializeN(svc.serial(INVOKE_SVC, bs.serialize(f, x), 500),
1)
    return out
end

function big(input)
    out = bs.deserializeN(svc.serial(BIG_SVC, bs.serialize(input), 1000), 1)
    return out
end

-- note that this function will ALWAYS generate a timeout if it
-- succeeds, because it shuts the slave down to execute the command:
function execute(input)
    out = bs.deserializeN(svc.serial(EXECUTE_SVC, bs.serialize(input),
1000), 1)
    return out
end

function quit()
    svc.serial(QUIT_SVC, "", 0)
    return nil
end
```

## remote.lua

```
print("remote ...")

-- load common definitions ...
common = dofile('common.lux')

-- this is used by ALOHA for serial services ...
port_0_index = {
    [INVOKE_SVC]      = "invoke",
    [BIG_SVC]         = "big",
    [EXECUTE_SVC]     = "execute",
}

-- this is used by ALOHA for WiFi services ...
rpc_network = {
    ["SSID"]          = "MyNetwork";      -- set to "" if propeller offers AP
    ["PASSPHRASE"]    = "TellMeASecret";   -- set to "" for a propeller AP

    [INVOKE_SVC]      = "xxx.xxx.xxx.xxx";
    [QUIT_SVC]        = "xxx.xxx.xxx.xxx";
    [BIG_SVC]         = "xxx.xxx.xxx.xxx";
    [EXECUTE_SVC]     = "xxx.xxx.xxx.xxx";
}

-- this is the function the server executes in the background ...
function background()
    -- the default is to do nothing!
end

-- define local services ...
-- this service quits the local server ...
function quit()
```

Catalina C

Aloha from Lua

```
print("Client requested shutdown\n")
os.exit()
end

print("... loaded")
```

## Appendix B - Life

This appendix contains details of a more complex eLua/ALOHA client/server demo.

The folder `demos/elua/life` contains two Lua implementations of Conway's Game of Life<sup>15</sup> - a simple cellular automaton.

First, the version of the program in `life.lua` is *not* a client/server program. It is a normal Lua program and it can be compiled and executed using normal Lua. For example:

```
luac -o life.lux life.lua
luax life.lux
```

When executed this way and using a serial HMI, each iteration of Life takes about 5 seconds. The best we can do on one Propeller is to recompile Lua to execute in NATIVE mode from Hub RAM. If we do that each iteration of Life takes about 1.6 seconds.

Another version of the Game of Life is in the files `common.lua`, `client.lua` and `server.lua` (and their compiled equivalents `common.lux`, `client.lux` and `server.lux`). This is an eLua client/server version of the same program. The server implements the cellular automaton, and the client handles the display of the output. Execute it using `eluax` with the command:

```
eluax client server
```

or just:

```
eluax
```

When executed this way, each iteration of Life takes about 5 seconds - i.e. about the same speed as the original version of the program - but now it is being executed from XMM, which allows for larger Life universes.

But with ALOHA, the client and server do not need to execute on the same propeller. To demonstrate this, two propellers are needed, prepared as a P2\_MASTER and P2\_SLAVE, as described in the main body of this document.

On the P2\_SLAVE<sup>16</sup>, execute `slave` but specify `remote` as the *client*.

```
slave remote server
```

Then, on the P2\_MASTER, execute `master` but specify `remote` as the *server*.

```
master client remote
```

---

<sup>15</sup> See [https://en.wikipedia.org/wiki/Conway's\\_Game\\_of\\_Life](https://en.wikipedia.org/wiki/Conway's_Game_of_Life)

<sup>16</sup> The commands *must* be executed in this order - i.e. the slave has to be started *before* the master, or the master will time out.

Using ALOHA this way doesn't speed the program up any further. But that's not the end of the story ...

On the P2\_SLAVE, we can use **sluafix** to execute the slave in NATIVE mode from Hub RAM:

**sluafix server**

or just:

**sluafix**

And on the P2\_MASTER, we can use **masterfx** to execute the client in NATIVE mode from Hub RAM:

**masterfx client remote**

When executed this way, each iteration of Life takes only 0.8s - twice as fast as the best that can be achieved on a single Propeller.

As Mr Spock might have said to Captain Kirk ...

*"It's Life, Jim, but not as we know it"*